

Chapter 5: Inference for numerical variables

Learning Objectives

Reading: Section 5.1 of OpenIntro Statistics

1. Define observations as paired if each observation in one dataset has a special correspondence or connection with exactly one observation in the other data set.
2. Carry out inference for paired data by first subtracting the paired observations from each other, and then treating the set of differences as a new numerical variable on which to do inference (such as a confidence interval or hypothesis test for the average difference).

Test yourself:

1. 20 cardiac patients' blood pressure is measured before taking a medication, and after. For a given patient, are the before and after blood pressure measurements dependent (paired) or independent?
2. A random sample of 100 students were obtained and then randomly assigned into two equal sized groups. One group went on a roller coaster while the other in a simulator at an amusement park. Afterwards their blood pressure measurements were taken. Are the measurements paired or not?

Reading: Section 5.2 of OpenIntro Statistics

3. Calculate the standard error of the difference between means of two independent samples as

$$SE = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}},$$

and use this standard error in hypothesis testing and confidence intervals comparing means of independent groups.

4. Recognize that a good interpretation of a confidence interval for the difference between two parameters includes a comparative statement.
5. Recognize that a confidence interval for the difference between two parameters that doesn't include 0 is in agreement with a hypothesis test where the null hypothesis that sets the two parameters equal to each other is rejected.

Test yourself:

1. Describe how the two sample means test is different from the paired means test, both conceptually and in terms of the calculation of the standard error.
2. A 95% confidence interval for the difference between the number of calories consumed by mature and juvenile cats ($\mu_{mat} - \mu_{juv}$) is (80 calories, 100 calories). Interpret this interval, and determine if it suggests a significant difference between the two means.

Reading: Sections 5.3 and 5.4 of OpenIntro Statistics

6. Explain why means of small samples ($n < 30$) follow the t -distribution under certain conditions (instead of the normal, z , distribution).
7. Describe how the t distribution is different from the normal distribution, and what "heavy tail" means in this context.

8. Note that the t distribution has a single parameter, degrees of freedom, and as the degrees of freedom increases, this distribution approaches the normal distribution.
9. Use a t statistic, with degrees of freedom $df = n - 1$ for inference for a population mean using data from a small sample:

$$\text{CI: } \bar{x} \pm t_{df}^* SE \qquad \text{HT: } T_{df} = \frac{\bar{x} - \mu}{SE}$$

where $SE = \frac{s}{\sqrt{n}}$.

10. Use a t statistic, with degrees of freedom $df = \min(n_1 - 1, n_2 - 1)$ for inference for difference between means of two population means using data from two small samples, where $SE = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$.
11. Make note of the pooled standard deviation but use it in very rare circumstances where the standard deviations of the populations being compared are known to be very similar.
12. Describe how to use the T table to obtain a p-value and a critical t-score (t_{df}^*) for a confidence interval.

Test yourself:

1. What is t^* for a 95% confidence interval for a mean, where the sample size is 13.
2. What is the p-value for a hypothesis test where the alternative hypothesis is two-sided, the sample size is 20, and the test statistic, T , is calculated to be 1.75?

Reading: Section 5.5 of OpenIntro Statistics

13. Define analysis of variance (ANOVA) as a statistical inference method that is used to determine if the variability in the sample means is so large that it seems unlikely to be from chance alone by simultaneously considering many groups at once.
14. Recognize that the null hypothesis in ANOVA sets all means equal to each other, and the alternative hypothesis suggest that at least one mean is different.

$$H_0 : \mu_1 = \mu_2 = \dots = \mu_k$$

H_A : At least one mean is different

15. List the conditions necessary for performing ANOVA

- (1) the observations should be independent within and across groups
- (2) the data within each group are nearly normal
- (3) the variability across the groups is about equal

and check if they are met using graphical diagnostics.

16. Recognize that the test statistic for ANOVA, the F statistic, is calculated as the ratio of the mean square between groups (MSG, variability between groups) and mean square error (MSE, variability within errors), and has two degrees of freedom, one for the numerator ($df_G = k - 1$, where k is the number of groups) and one for the denominator ($df_E = n - k$, where n is the total sample size).

- The expectation is not to compute MSG or MSE by hand, which in practice is done using statistical software, but rather to have a conceptual understanding of how they are calculated and what they measure.

17. Describe why calculation of the p-value for ANOVA only considers the upper tail of the F distribution and never the lower tail.

18. Describe why conducting many t tests for differences between each pair of means leads to an increased Type 1 Error rate, and we use a corrected significance level (Bonferroni corection, $\alpha^* = \alpha/K$, where K is the number of comparisons being considered) to combat inflating this error rate.
19. Describe why it is possible to reject the null hypothesis in ANOVA but not find significant differences between groups as a result of pairwise comparisons.

Test yourself:

1. We would like to compare the average income of Americans who live in the Northeast, Midwest, South, and West. What are the appropriate hypotheses?
2. Suppose the sample in the question above has 1000 observations, what are the degrees of freedom associated with the F-statistic?
3. Suppose the null hypothesis is rejected. Describe how we would discover which regions' averages are different from each other? Make sure to discuss how many pairwise comparisons we would need to make, and what the corrected significance level would be.
4. What visualizations are useful for checking each of the conditions required for performing ANOVA?